

# When to Trust the Peg

Using Random Matrix Theory to Validate  
the Anchored Pricing Assumption in Real Time

*CZ*

May 2025

## Abstract

The LlamaRisk October 2025 report identified a critical oracle design failure: Binance’s internal order book pricing triggered procyclical liquidation cascades under stress, while Aave’s USDT-anchored pricing for USDe correctly isolated venue-specific price distortion. However, the report relied on an implicit assumption never formally validated in real time—is the peg actually valid, or is it concealing genuine solvency risk? This paper applies Random Matrix Theory (RMT) to real-time cross-stablecoin correlation analysis, proposing the ratio of the maximum eigenvalue to the Marchenko-Pastur upper bound ( $\lambda$ -ratio) as a single quantifiable verification signal. Core empirical findings:  $\lambda$ -ratio never breached 1.0 across a 45-day baseline period (zero false positives); it first breached 1.0 at 12:00 UTC on October 10, 2025—**9 hours and 36 minutes** before Binance’s officially confirmed depeg onset at 21:36 UTC; and it remained elevated for two weeks before normalising. This breach-then-recovery pattern is structurally distinct from a solvency crisis signature, and retrospectively validates Aave’s decision to maintain anchored pricing. This paper also reports a null result on eigenvector loading direction as a precursor signal: in the noise regime the principal eigenvector flips simultaneously for all assets due to numerical instability, producing no asset-specific information. Reporting null results is treated as equally important to reporting positive findings.

**Keywords:** oracle design, stablecoin depeg, random matrix theory, Marchenko-Pastur law, DeFi risk, USDe, Aave, LlamaGuard

# Contents

|          |                                                                      |           |
|----------|----------------------------------------------------------------------|-----------|
| <b>1</b> | <b>Background and Open Question</b>                                  | <b>3</b>  |
| 1.1      | What the LlamaRisk Report Solved . . . . .                           | 3         |
| 1.2      | Why Anchored Pricing Could Work Here . . . . .                       | 3         |
| 1.3      | The Open Question . . . . .                                          | 4         |
| <b>2</b> | <b>The Signal: <math>\lambda</math>-Ratio</b>                        | <b>4</b>  |
| 2.1      | The Marchenko-Pastur Law . . . . .                                   | 4         |
| 2.2      | Definition of $\lambda$ -Ratio . . . . .                             | 4         |
| 2.3      | Effective Rank as a Secondary Diagnostic . . . . .                   | 5         |
| <b>3</b> | <b>Empirical Results</b>                                             | <b>5</b>  |
| 3.1      | Data and Methodology . . . . .                                       | 5         |
| 3.2      | 45-Day Baseline: Zero False Positives . . . . .                      | 5         |
| 3.3      | Signal Timeline and Lead Time . . . . .                              | 6         |
| 3.4      | Interpreting the Breach-Then-Recovery Pattern . . . . .              | 6         |
| <b>4</b> | <b>Eigenvector Loading Analysis: A Complementary Diagnostic Tool</b> | <b>7</b>  |
| 4.1      | The Test . . . . .                                                   | 7         |
| 4.2      | The Result . . . . .                                                 | 8         |
| 4.3      | Interpretation and Implications . . . . .                            | 8         |
| <b>5</b> | <b>Recommendation for the LlamaGuard Verification Layer</b>          | <b>9</b>  |
| <b>6</b> | <b>Limitations</b>                                                   | <b>10</b> |
| <b>7</b> | <b>Conclusion</b>                                                    | <b>10</b> |


# 1 Background and Open Question

## 1.1 What the LlamaRisk Report Solved

The LlamaRisk October 2025 report [LlamaRisk, 2025] makes two concrete contributions to DeFi oracle design.

First, it documents a *venue-specific price distortion*: during the October 10 USDe depeg event, Binance’s internal order book reported a price as low as \$0.62 while prices on other venues remained close to \$1.00 (Figure 1). This is not a market-wide pricing failure; it is a localised order book microstructure event confined to a single venue.

Second, it traces the cascade mechanism: protocols using Binance’s internal pricing for collateral valuation triggered liquidations on this distorted price, creating further selling pressure, triggering further liquidations. Aave’s anchored pricing—pegging USDe to USDT rather than using real-time spot prices—interrupted this feedback loop by refusing to transmit the Binance distortion into collateral valuation.



**Figure 1:** USDe price during the October 10–11, 2025 event window. The red line shows the CoinGecko aggregated price series. The grey dashed line marks the \$1.00 peg. The orange dotted vertical line marks the official Binance depeg onset (21:36 UTC, October 10). Note that the aggregated price does not capture the \$0.62 Binance internal print; the venue-specific distortion documented in LlamaRisk [2025] was isolated to a single order book.

## 1.2 Why Anchored Pricing Could Work Here

Understanding why anchored pricing was the correct mechanism requires distinguishing two questions that are easy to conflate.

The Binance \$0.62 print was a *failure of price discovery*—not a reflection of USDe’s intrinsic value. Order book liquidity had dried up, market makers had withdrawn, and a large sell order moved price to an extreme level not reflected on other venues. Ethena’s protocol had no corresponding collapse in collateral value.

Aave correctly diagnosed this as venue-specific distortion and refused to transmit it to collateral valuation. That judgment was retrospectively correct: USDe’s price recovered, and the protocol suffered no bad debt.

However, anchored pricing can *only* isolate price discovery failures. It cannot and should not isolate genuine solvency failures. If USDe’s collateral had actually deteriorated—if the delta-

neutral hedging positions had failed at scale—then anchored pricing would have accumulated bad debt while hiding the risk.

### 1.3 The Open Question

**The methodological gap.** Aave’s decision to maintain anchored pricing was correct. But at the time the decision was made, there was no real-time mechanism to verify that the premise underlying the decision—“USDe has no solvency problem”—was actually true. The protocol was right, but not because a tool told it so.

The report proposes a LlamaGuard dynamic data feed with a verification layer [LlamaRisk, 2025], but does not specify what signals that layer should monitor, or how to distinguish “the peg is temporarily dislocated by liquidity noise” from “the peg is concealing a genuine solvency event” in real time. This paper answers that question.

## 2 The Signal: $\lambda$ -Ratio

### 2.1 The Marchenko-Pastur Law

Let  $\mathbf{X}$  be a  $T \times N$  matrix of standardised log-returns, where  $T$  is the number of time observations and  $N$  is the number of assets. The sample correlation matrix is  $\mathbf{C} = \frac{1}{T} \mathbf{X}^\top \mathbf{X}$ .

In the null hypothesis where returns are i.i.d. (no genuine correlation structure), Marchenko and Pastur [1967] proved that as  $T, N \rightarrow \infty$  with  $q = N/T$  fixed, the eigenvalue density of  $\mathbf{C}$  converges to a known distribution with support  $[\lambda_-, \lambda_+]$ :

$$\lambda_{\pm} = (1 \pm \sqrt{q})^2, \quad q = \frac{N}{T}. \quad (1)$$

Any eigenvalue exceeding  $\lambda_+$  cannot be explained by sampling noise alone; it represents a genuine common factor driving co-movement across assets [Laloux et al., 1999, Plerou et al., 1999]. This boundary is a theoretical result, not a fitted parameter.

### 2.2 Definition of $\lambda$ -Ratio

Define the  $\lambda$ -ratio as:

$$\lambda\text{-ratio} = \frac{\lambda_1}{(1 + \sqrt{q})^2}, \quad (2)$$

where  $\lambda_1$  is the largest observed eigenvalue. The interpretation is direct:

- $\lambda\text{-ratio} < 1.0$ : The correlation structure is consistent with pure noise. No genuine common factor is present. Price deviations are likely venue-specific distortions. The peg assumption is supported.

- $\lambda$ -ratio  $> 1.0$ : A genuine common factor has entered the system. Cross-stablecoin co-movement exceeds what noise alone can explain. The peg assumption requires active verification, not passive acceptance.

### 2.3 Effective Rank as a Secondary Diagnostic

The effective rank of the correlation matrix, introduced by [Roy and Vetterli \[2007\]](#), is defined as the exponential of the Shannon entropy of the normalised eigenvalue distribution:

$$R_{\text{eff}} = \exp\left(-\sum_i p_i \ln p_i\right), \quad p_i = \frac{\lambda_i}{\sum_j \lambda_j}. \quad (3)$$

This ranges from 1 (all variance concentrated in a single factor) to  $N$  (uniform distribution, maximum independence). A declining effective rank indicates the stablecoin basket is increasingly driven by a single dominant mode. The recovery of effective rank after a stress event provides a diagnostic signal: in genuine insolvency the common factor intensifies; in transient stress it dissipates.

For a comprehensive treatment of RMT in financial applications, see [Bouchaud and Potters \[2011\]](#).

## 3 Empirical Results

### 3.1 Data and Methodology

The analysis uses OHLC data for five stablecoins (USDe, USDT, USDC, DAI, FRAX) obtained via the CoinGecko public API, covering a 90-day window centred on the October 2025 event. Log-returns are computed as  $r_t = \ln(P_t/P_{t-1})$ . The rolling RMT analysis uses a window of 72 observations with a step of 4 observations. With  $N = 5$  and  $T = 72$ , we have  $q = 5/72 \approx 0.069$ , giving a Marchenko-Pastur upper bound of  $\lambda_+ \approx 1.547$  for the raw maximum eigenvalue, corresponding to  $\lambda$ -ratio = 1.0.

**Ground-truth benchmark.** Binance’s official post-incident statement confirms that the severe depeg began at **21:36 UTC on October 10, 2025**, with prices reaching their lowest point at 21:20–21:21 UTC.

### 3.2 45-Day Baseline: Zero False Positives

From September 15 through October 9—45 consecutive days—the  $\lambda$ -ratio ranged between 0.84 and 0.99. It never once breached 1.0. Effective rank was stable between 4.70 and 4.86 throughout this period (Figure 2).

A signal that fires frequently in calm periods is not useful for governance decisions. The Marchenko-Pastur bound, being a theoretical result rather than a fitted threshold, produces


**zero false positives** across the entire 45-day pre-event period without any parameter tuning.

### 3.3 Signal Timeline and Lead Time

Table 1 presents the  $\lambda$ -ratio and effective rank at key timestamps around the event. Figure 2 shows the full time series.

**Table 1:**  $\lambda$ -ratio and effective rank at key timestamps. The official depeg onset time is taken from Binance’s post-incident statement [LlamaRisk, 2025].

| Time (UTC)                    | $\lambda$ -ratio | Eff. Rank    | Status                                                |
|-------------------------------|------------------|--------------|-------------------------------------------------------|
| Sep 15 – Oct 9 (full period)  | 0.84–0.99        | 4.70–4.86    | Below MP bound. Zero breaches.                        |
| Oct 10, 04:00                 | 0.941            | 4.840        | Still in noise regime.                                |
| <b>Oct 10, 12:00</b>          | <b>1.013</b>     | <b>4.691</b> | <b>First breach. Signal fires.</b>                    |
| <i>Binance official depeg</i> | —                | —            | <i>21:36 UTC <math>\leftarrow</math> 9h 36m later</i> |
| Oct 11, 04:00                 | 1.226            | 4.337        | Confirmed structural break.                           |
| Oct 11–25                     | 1.20–1.25        | 4.29–4.44    | Sustained elevation. Two weeks.                       |
| Oct 26 onwards                | <0.85            | >4.86        | Full normalisation.                                   |



**Figure 2:**  $\lambda$ -ratio across the 90-day sample period. The red dashed line marks the Marchenko-Pastur upper bound ( $\lambda$ -ratio = 1.0). The red dash-dotted line marks the escalation threshold (1.2). The orange dotted vertical line marks the official depeg onset (21:36 UTC, October 10). The blue dashed vertical line marks the first signal breach (12:00 UTC, October 10). The shaded region indicates the period above the MP bound. Note zero breaches across the 45-day pre-event baseline.

### 3.4 Interpreting the Breach-Then-Recovery Pattern


Three features of the  $\lambda$ -ratio trajectory are diagnostically significant.

**Pre-breach signal.** The  $\lambda$ -ratio crossed 1.0 at 12:00 UTC while USDe prices had not yet shown significant dislocation. The cross-stablecoin correlation structure encoded stress before the price signal was visible—9 hours and 36 minutes before official confirmation.

**Two-week persistence.** A genuine liquidity event resolves within 24–48 hours as arbitrageurs

close the gap and order books rebalance. The  $\lambda$ -ratio remained above 1.20 for two full weeks, indicating a lasting structural change in the correlation regime, not a transient fluctuation.

**Recovery.** The  $\lambda$ -ratio fell back below 0.85, and effective rank recovered to its highest level in the entire dataset—4.86–4.92, above pre-event levels (Figure 3). This recovery trajectory is structurally incompatible with a solvency crisis. In genuine insolvency, the common factor that enters the system does not subsequently dissipate; it persists and intensifies as bad debt accumulates.



**Figure 3:** Effective rank of the cross-stablecoin correlation matrix across the 90-day sample period. The red dashed line marks the theoretical minimum ( $R_{\text{eff}} = 1$ , complete collapse). The orange dotted and blue dashed vertical lines mark the official depeg and signal breach timestamps respectively (see Figure 2). Note that effective rank recovers above its pre-event level after October 26, consistent with risk-off rebalancing and inconsistent with a solvency crisis trajectory.

**Key finding.** The breach-then-recovery pattern retrospectively validates Aave’s decision to maintain anchored pricing. The peg premise—that USDe had no solvency problem—was correct. The  $\lambda$ -ratio, had it been monitored in real time, would have provided quantitative support for this conclusion throughout the event: below 1.0 before the event (maintain peg), above 1.0 during the event (verify actively), and returning below 1.0 afterwards (peg premise revalidated).

## 4 Eigenvector Loading Analysis: A Complementary Diagnostic Tool

This section examines whether eigenvector loadings can provide additional insight beyond the primary  $\lambda$ -ratio signal. While the initial hypothesis of using loading direction/sign as a precursor signal did not hold, a refined Leave-One-Out approach yields valuable post-event diagnostic information.

### 4.1 The Test

For each rolling window, the dominant eigenvector (corresponding to  $\lambda_1$ ) was computed using a **Leave-One-Out** procedure: the principal eigenvector is calculated on all assets except the


one under consideration, and the excluded asset is then projected onto this “pure benchmark” to measure its idiosyncratic loading.

## 4.2 The Result

As shown in Figure 4, the Leave-One-Out projected loadings reveal two distinct regimes:

- **Pre-event period (Sep 15 – Oct 10, 2025):** All five stablecoins exhibit relatively synchronized and moderate fluctuations, with no clear divergence observable. This is consistent with the noise regime identified by  $\lambda$ -ratio  $< 1.0$ .
- **Post-event period (after Oct 10):** Only USDe shows extreme and persistent flickering, with its projected loadings oscillating violently between approximately  $+0.6$  and  $-0.6$ . In contrast, USDT, USDC, DAI, and FRAX remain comparatively stable, mostly fluctuating within  $\pm 0.2$ .

This demonstrates that the Leave-One-Out method successfully isolates USDe’s idiosyncratic behavior following the depeg event.



**Figure 4:** Leave-One-Out Projected Loadings on the dominant eigenvector. Only USDe (red solid line) exhibits extreme post-event flickering, while the other stablecoins remain relatively stable. The orange dotted vertical line marks the USDe depeg event on October 10, 2025.

## 4.3 Interpretation and Implications

The Leave-One-Out projected loading analysis does **not** provide a reliable precursor signal before  $\lambda$ -ratio breaches 1.0. However, it serves as a strong **diagnostic** tool once stress has been detected:

- It confirms that USDe became the primary source of idiosyncratic stress after the event, rather than simply participating in broad systemic contagion.
- It forms an effective complement to  $\lambda$ -ratio:
  - $\lambda$ -ratio acts as an **early warning** signal (breaching  $\sim 9.6$  hours before the official announcement);

- Leave-One-Out projected loadings provide **post-event confirmation and attribution**, identifying USDe as the decoupled asset.

This contrast reinforces a key theoretical insight: in the noise regime ( $\lambda$ -ratio  $< 1.0$ ), eigenvector directions are numerically unstable. Eigenvector analysis only becomes informative after a genuine common factor emerges ( $\lambda$ -ratio  $> 1.0$ ).

The finding supports the hierarchical decision framework in Section 5: use  $\lambda$ -ratio as the primary gatekeeper, followed by more granular tools such as Leave-One-Out loadings for diagnosis and attribution.

## 5 Recommendation for the LlamaGuard Verification Layer

The original report proposes a three-layer LlamaGuard architecture (valuation, verification, execution) but does not specify what the verification layer should monitor [LlamaRisk, 2025]. The  $\lambda$ -ratio provides that specification.

**Table 2:** Decision framework for the LlamaGuard verification layer based on  $\lambda$ -ratio. The 1.0 threshold is theoretically derived; the 1.2 escalation level and 48-hour persistence criterion are calibrated to this event and require cross-validation (Section 6).

| $\lambda$ -ratio                          | Interpretation                                                       | Protocol Action                                                                                                                                     |
|-------------------------------------------|----------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------|
| $< 1.0$                                   | Noise regime. No genuine common factor.                              | <b>Maintain anchored pricing.</b> Do not transmit venue distortion to collateral valuation. ( <i>Zero false positives across 45-day baseline.</i> ) |
| <b>Crosses 1.0</b>                        | Genuine common factor confirmed. Peg premise needs verification.     | <b>Trigger governance review.</b> Do not immediately change pricing mechanism. Monitor for persistence.                                             |
| $> 1.2$ , <b>sustained</b> $> 48\text{h}$ | Structural stress confirmed. Not resolving as liquidity noise would. | <b>Escalate.</b> Evaluate solvency exposure. Consider suspending new borrowing.                                                                     |
| <b>Returns to</b> $< 1.0$                 | Breach-then-recovery. Solvency crisis ruled out.                     | <b>Restore normal operations.</b> Peg premise revalidated.                                                                                          |

**Implementation note.** The  $\lambda$ -ratio requires no new data sources—only the cross-stablecoin price data the valuation layer already collects. The eigenvalue decomposition of a  $5 \times 5$  matrix is computationally negligible and can be updated in real time at any price feed frequency.

## 6 Limitations

**Single-event calibration.** The  $\lambda$ -ratio = 1.0 threshold is theoretically motivated and requires no tuning. However, the escalation threshold of 1.2 and the 48-hour persistence criterion are calibrated to this one event. Cross-validation against other stablecoin stress episodes—the USDC depeg in March 2023, FRAX rebalancing events, crvUSD emergency auctions—is required before these secondary thresholds are used in production.

**Finite-sample corrections.** With  $N = 5$  and  $T = 72$ , the Marchenko-Pastur law is asymptotically valid but finite-sample deviations may be material. Tracy-Widom corrections or bootstrap-calibrated thresholds should be explored for small- $N$  applications.

**Causality.**  $\lambda$ -ratio detects that a genuine common factor is present; it does not identify what that factor is. A breach could reflect USDe-specific solvency stress, correlated macro shocks, or microstructure regime changes unrelated to any single asset. Human judgment remains essential—the signal is a gate, not a verdict.

**Asset set dependency.** The correlation matrix is computed across a fixed basket of five stablecoins. Adding or removing assets changes the matrix dimensions and therefore the MP upper bound, requiring recalibration.

## 7 Conclusion

Aave’s anchored pricing was the right decision for October 2025. But it was right without a real-time tool to verify the premise it depended on. This paper proposes  $\lambda$ -ratio as that tool.

Its properties are well-suited to the task: the 1.0 threshold is derived from the Marchenko-Pastur law [Marchenko and Pastur, 1967], not empirically fitted; zero false positives across 45 baseline days; 9 hours and 36 minutes of lead time before official depeg confirmation; and a breach-then-recovery trajectory that distinguishes structural stress from solvency failure, consistent with the effective rank framework of Roy and Vetterli [2007].

The null result on eigenvector loading direction is reported with equal prominence. The loading sign is numerically unstable in the noise regime where a precursor would be most useful—making it unsuitable as a standalone signal. This is not a failure of the RMT framework; it is a clarification of which component carries the information.

**The decision rule:** trust the peg when  $\lambda$ -ratio  $< 1.0$ ; verify the peg premise actively when it breaches 1.0; treat the return below 1.0 as confirmation that the solvency premise is intact. Applied to October 2025, this rule would have supported maintaining anchored pricing throughout—and would have told the protocol, in real time, why.

## References

- Jean-Philippe Bouchaud and Marc Potters. Financial applications of random matrix theory: A short review. In Gernot Akemann, Jinho Baik, and Philippe Di Francesco, editors, *The Oxford Handbook of Random Matrix Theory*, pages 824–850. Oxford University Press, 2011.
- Laurent Laloux, Pierre Cizeau, Jean-Philippe Bouchaud, and Marc Potters. Noise dressing of financial correlation matrices. *Physical Review Letters*, 83(7):1467–1470, 1999. doi: 10.1103/PhysRevLett.83.1467.
- LlamaRisk. When pricing breaks: USDe. Technical report, LlamaRisk Research, October 2025. URL <https://llamarisk.com/research/when-pricing-breaks-usde>. Accessed May 2025.
- Vladimir A. Marchenko and Leonid A. Pastur. Distribution of eigenvalues for some sets of random matrices. *Mathematics of the USSR-Sbornik*, 1(4):457–483, 1967. doi: 10.1070/SM1967v001n04ABEH001994.
- Vasiliki Plerou, Parameswaran Gopikrishnan, Bernd Rosenow, Luís A. N. Amaral, and H. Eugene Stanley. Universal and nonuniversal properties of cross correlations in financial time series. *Physical Review Letters*, 83(7):1471–1474, 1999. doi: 10.1103/PhysRevLett.83.1471.
- Olivier Roy and Martin Vetterli. The effective rank: A measure of effective dimensionality. In *Proceedings of the 15th European Signal Processing Conference (EUSIPCO)*, pages 606–610, 2007.